

# Concurrent Programming (Part II)

## Lecture 11: Concurrency without access to shared memory (i.e. distributed systems)

Dr Kevin Bryson

[K.Bryson@cs.ucl.ac.uk](mailto:K.Bryson@cs.ucl.ac.uk)


Room 8.04

Course Web Site on Moodle

<http://moodle.ucl.ac.uk/course/view.php?id=753>

Enrolment Key: ATOMIC


## Recall shared memory architectures from Lecture 2 ...



## **Thread communication within a single JVM (within one OS process) using shared memory**

- Until now all our programs have assumed threads to be in the same JVM and therefore sharing an address space
- How do these threads communicate?
  - Reading/writing shared variables (generally under restricted concurrent access such as mutual exclusion)
  - Event conditions (wait/notify mechanism)
  - Changing the state of more complex data structures generally structured as Monitors (for instance a buffer)
  - Directly calling thread methods such as `interrupt()`

Also recall we briefly showed a distributed system in Lecture 2 ... what are the key concepts for understanding concurrency in such systems?



## Well ... what is the relevance of this picture?

- This picture begs the question ... how do we communicate with each other? (Given that we do not share direct access to each other's working memories!)
- Distributed threads need to employ a similar mechanisms ...



## Overview of this lecture

- When processes do not directly share an address space they cannot communicate by means of shared local variables like all the examples we have had so far!
- One example would be when processes sit at different ends of a network (i.e. on different physical machines).
- This lecture is about the *key theoretical abstractions* behind concurrent systems that do not employ shared memory mechanisms.
- This will lead into the final set of lectures that will cover more concrete material about writing distributed concurrent systems using Java.

## Key concept 1 - Message Passing

- Messages can be sent to communicate from one thread to another Thread. These threads may be within:
  - the same OS process/JVM (as before ...)
  - different OS processes on the same machine
  - different OS processes on different physical machines, with a communication network between them
- A sender process **sends** the message and a receiver process **receives** the message
- BUT WHY WOULD WE WANT TO SEND MESSAGES BETWEEN THREADS IN THE SAME JVM ???

## Key concept 2 – concurrency behaviour

- **Synchronous** message passing: the sender sends the message and then *blocks* – waiting for the receiver to accept it. The sender only continues once the message has been received.
- **Asynchronous** message passing: the sender sends the message and then continues executing. This is a *non-blocking* protocol.



## Key concept 3 – communication path

- **One-to-one over a dedicated ‘network’.** A dedicated channel between the sender and receiver exists (for instance in Occam where the physical transputer chips are connected in a North/South/East/West fashion to other transputer chips *in hardware*).
- **One-to-one over a switching ‘network’.** In this case the receiver processes need unique identifiers so that messages can be correctly routed to them.
- **One-to-many messaging (Broadcast messaging).** The sender does not need an identifier for the receiver since the message gets sent to all processes currently listening.
- **THESE ARE ABSTRACT CONCEPTS – communication could be between threads in the same JVM using a message passing package like Java Message Service (JMS), i.e. in software without a physical network. The above concepts still apply!**


## Key concept 4 – messaging protocols

- **One-way communication.** A communication event only consists of the sender sending a message to the receiver.
- **Two-way communication: *Rendezvous* (or request-reply):** the sender sends the message and then *blocks* while it waits for the receiver to send back a message back. The receiver typically carries out processing on the message before replying.

## Examples of message passing

- When thinking about examples we need to identify the ‘processes’ doing the communication and what constitutes the ‘message’ (i.e. the level of abstraction)
- There are different ways in which messages can be sent.  
Can you think of some from real life?  
Can you think of key differences in their mechanisms?
- **Synchronous:**
  - fax -> fax (message on bit of paper)
  - person -> person on phone
  - person -> answer-phone (message what is said)
- **Asynchronous:**
  - person -> person via a letter, email, ...
  - person -> person **via answer-phone**
- **Rendezvous:**
  - person -> person (telephone call)
  - person -> person (meeting and chatting ...)

## 10.1 Synchronous Message Passing - **channel**



♦ **send(e, c)** - send the value of the expression `e` to channel `c`. The process calling the send operation is **blocked** until the message is received from the channel.

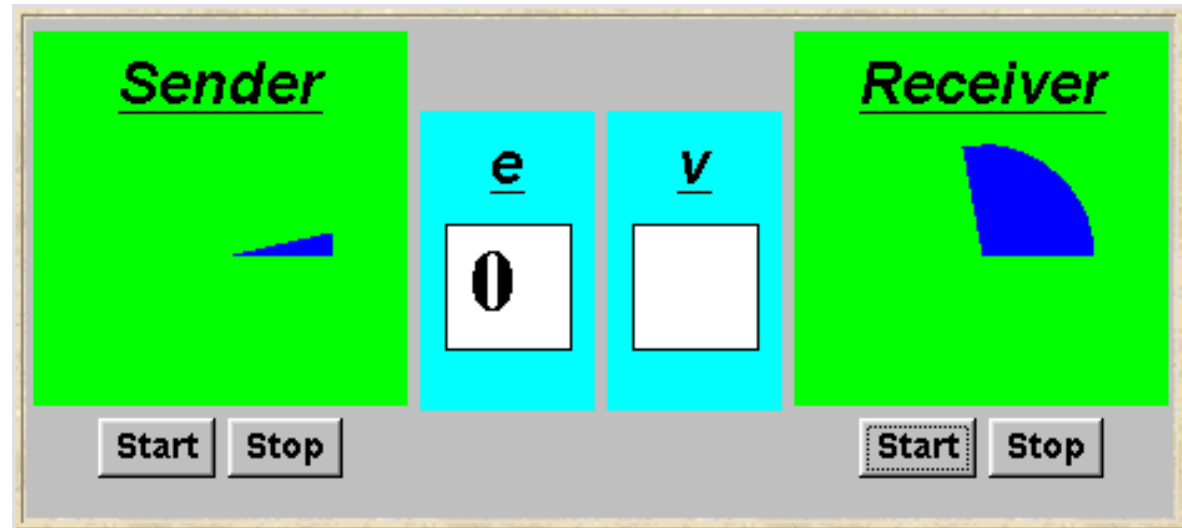
♦ **v = receive(c)** - receive a value into local variable `v` from channel `c`. The process calling the receive operation is **blocked** waiting until a message is sent to the channel.

# Synchronous message passing - demo

A sender communicates with a receiver using a single **channel**.

The sender sends a sequence of integer values from 0 to 9 and then restarts at 0 again.

See the demo ...



```
Channel<Integer> chan = new Channel<Integer>();  
tx.start(new Sender(chan, senddisp));  
rx.start(new Receiver(chan, recvdisp));
```


## But does this not just implement the Channel as a monitor using shared variables ?

- **YES !**
- ... but it is still message passing since it is the approach that is important ... not the implementation (it is an abstract concept).
- Messages are written by a sender process ... and then they are received by a receiver process ... the 'ownership' of the message is transferred.  
(The Java Messaging Service (JMS API) uses shared memory when it is possible/efficient to do so ...)
- **In theory a hardware communication channel could have existed between these processes ...**

# Car Park Demo – key difference between shared memory & message passing versions


## Shared Memory

Calling Methods



## Message Passing

Transmitting Messages – could be distributed



**PROBLEM - WHAT HAPPENS IF THE CAR PARK CONTROLLER HAD BLOCKED WAITING ON THE 'WRONG' MESSAGE ???**


## Selective Receives

- What happens if a process wants to receive from more than one channel?
  - It blocks if it chooses the wrong one!
- Selective receive is a mechanism employed in languages such as Occam and Ada to solve this problem.
- The mechanism is similar to a 'select' statement which blocks until it can receive data on a particular channel.
- G1, G2, G3 are Boolean *guards* which must be true for the receive to be eligible.
- S1, S2, S3 are statements which are carried out depending on which channel it receives data on.

```
select  
    when G1 and v1 = receive(chan1) => S1;  
    when G2 and v2 = receive(chan1) => S2;  
    when G3 and v3 = receive(chan1) => S3;  
end
```



## 10.2 Asynchronous Message Passing - port



♦ **send(e, p)** - send the value of the expression  $e$  to port  $p$ . The process calling the send operation is **not blocked**. The message is queued at the port if the receiver is not waiting.

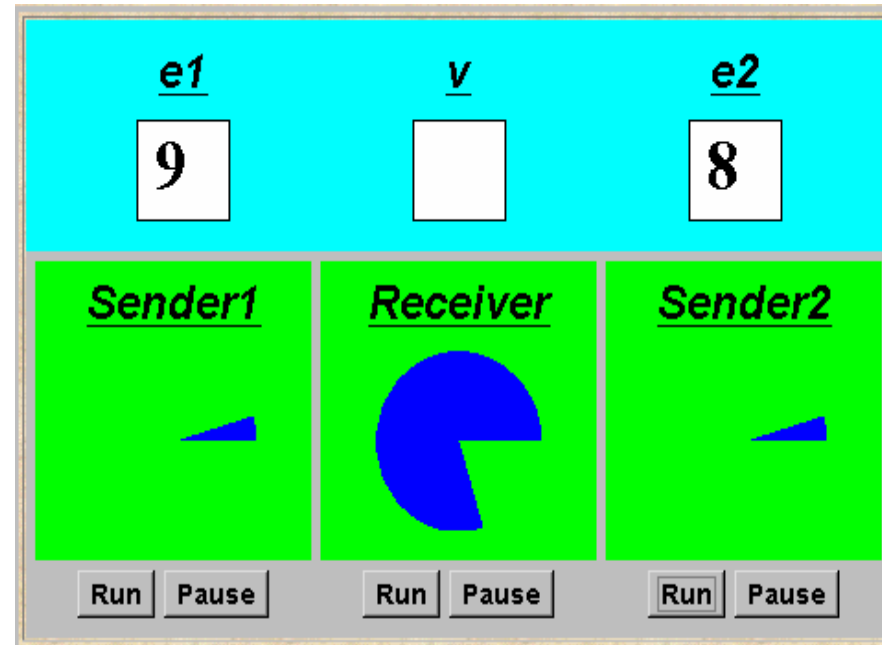
♦ **v = receive(p)** - receive a value into local variable  $v$  from port  $p$ . The process calling the receive operation is **blocked** if there are no messages queued to the port.

## Asynchronous message passing - demo

Two senders communicate with a receiver via an “unbounded” **port**.

Each sender sends a sequence of integer values from 0 to 9 and then restarts at 0 again.

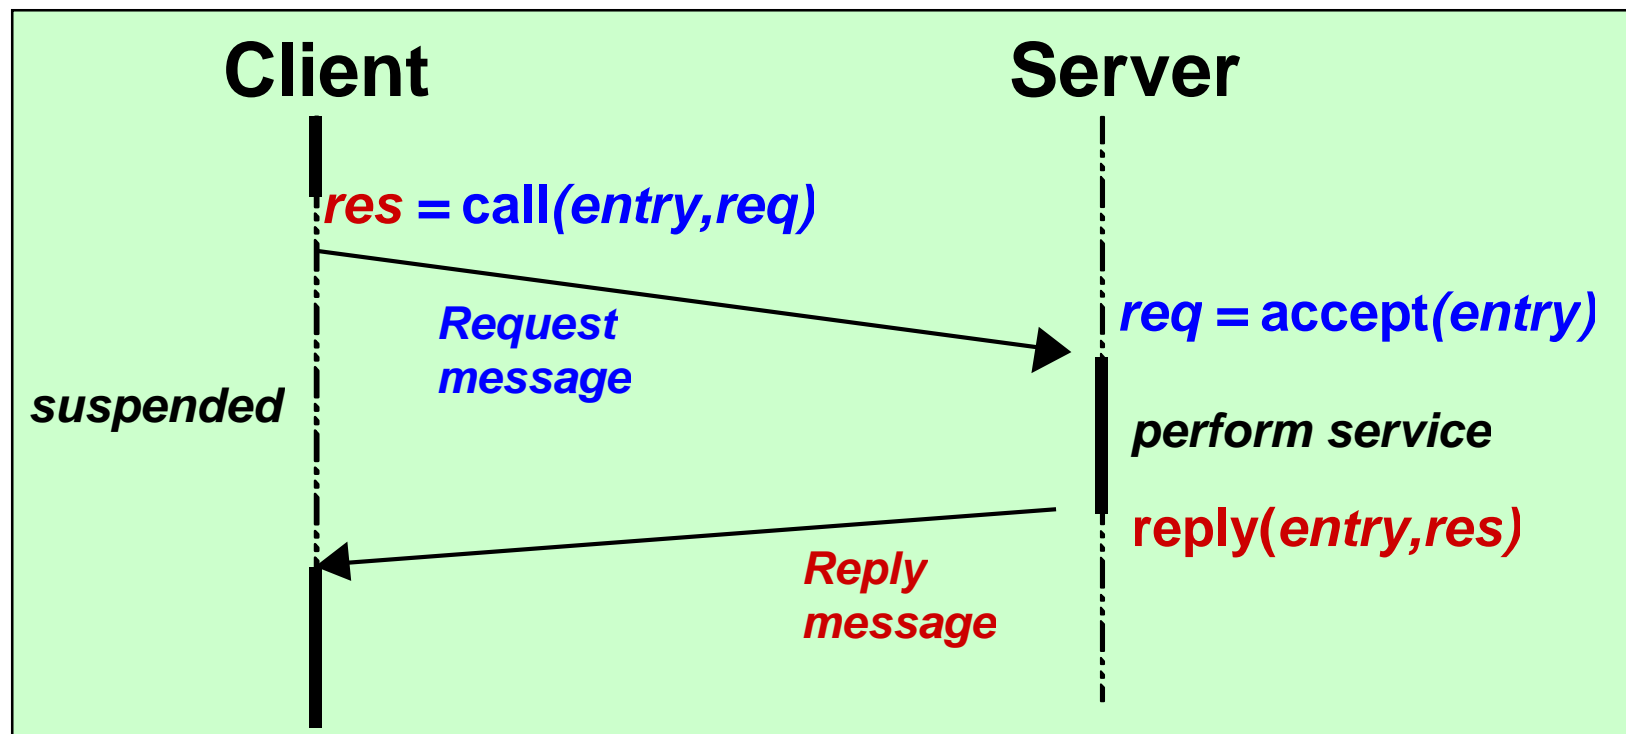
**See the demo ...**



```
Port<Integer> port = new Port<Integer> ();  
tx1.start(new Asender(port,send1disp));  
tx2.start(new Asender(port,send2disp));  
rx.start(new Areceiver(port,recvdisp));
```

## Rendezvous

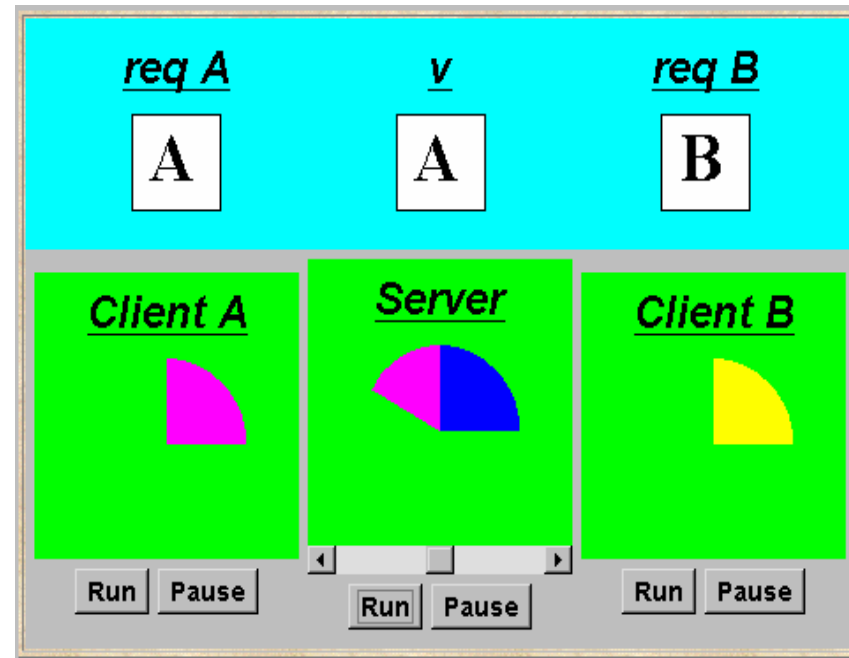
Rendezvous is a form of **request-reply** to support **client-server** communication. Many clients may request service, but only one is serviced at a time.



## Rendezvous - demo

Two clients call a server which services a request at a time.

See the demo ...



```
Entry<String,String> entry = new Entry<String,String> ();  
clA.start(new Client(entry,clientAdisp,"A"));  
clB.start(new Client(entry,clientBdisp,"B"));  
sv.start(new Server(entry,serverdisp));
```

## Summary

- When threads are not in the same address space they cannot communicate using shared memory
- They can communicate using message passing
- Message passing *can be* used even when threads do share a common address space ... to facilitate 'de-coupling' of the components. But it *is essential* when threads/processes are physically distributed.
- We have seen different types of message passing:
  - Synchronous
  - Asynchronous
  - Rendezvous
- **Background reading** – Magee & Kramer, Chapter 10.