

Concurrent Programming (Part II)

Lecture 14: Multi-Threaded Socket-Based Servers and Thread Pools

Dr Kevin Bryson

K.Bryson@cs.ucl.ac.uk

Room 8.04

Course Web Site on Moodle

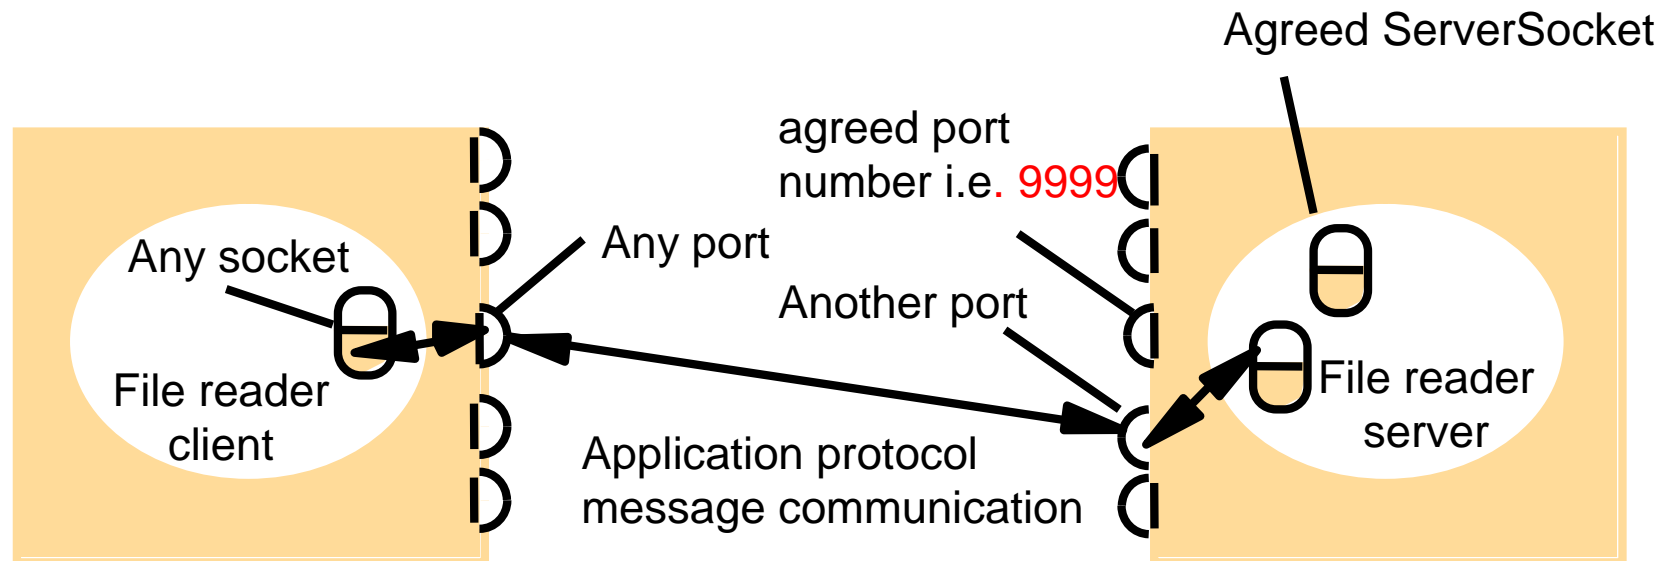
<http://moodle.ucl.ac.uk/course/view.php?id=753>

Enrolment Key: ATOMIC

Overview

- In the last lecture we looked at how to construct a client/server architecture using sockets.
- This was a single-threaded server, it could only service one client after another client (although buffering of the socket streams gave the illusion it may be multi-threaded ...)
- In this lecture we will examine how to make the server multi-threaded and how server response time is affected.
- Key Question: What factors affect server performance ?
- This question will lead us into a discussion of Thread Pools.

Recall the File Reader Server Demo



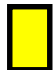


- Recall that the FileReaderServer is single threaded.
- The single 'Main' thread blocks waiting for a request at the ServerSocket accept() method.
- It accepts a client connection and continues to process it.
- Only when it is finished will it accept another connection.

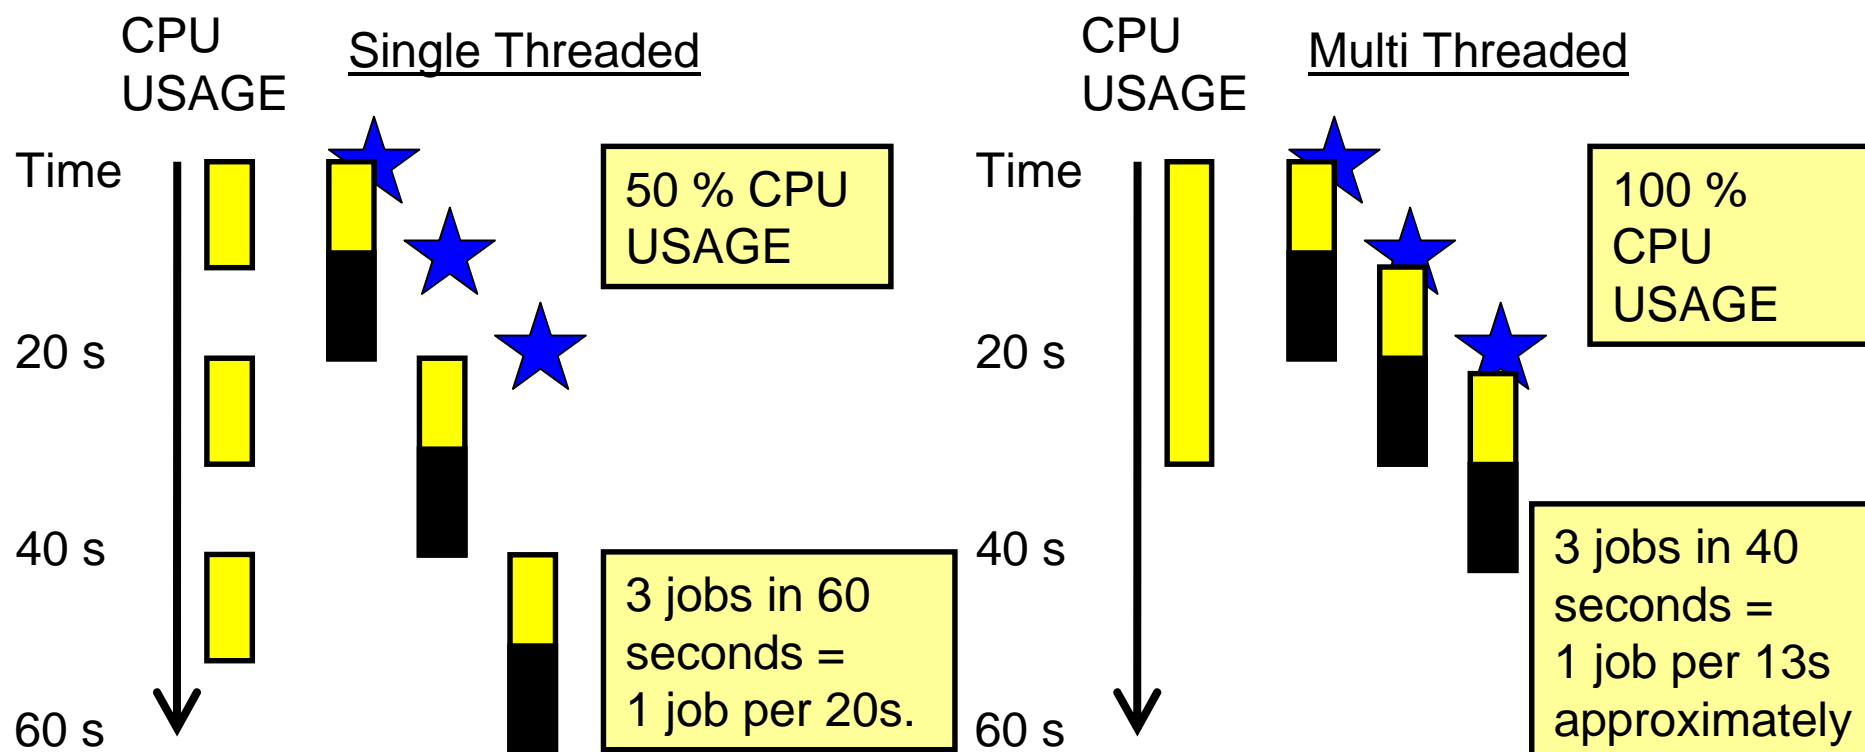
Why would we want to make the server multi-threaded? Discuss ...

- If tasks **block** (for instance on I/O), even a single processor will not be optimally used (since it will be idle when the Thread blocks).
- If the server has multiple processors – these may not be fully utilized executing a single task at a time
- In these cases the **throughput** will be lower than ideal (i.e. the average number of jobs processed per unit time)
- ***Making the server multi-threaded can increase processor utilization and hence throughput (although not always!)***

Tasks blocking on IO - increasing throughput

Consider this scenario:

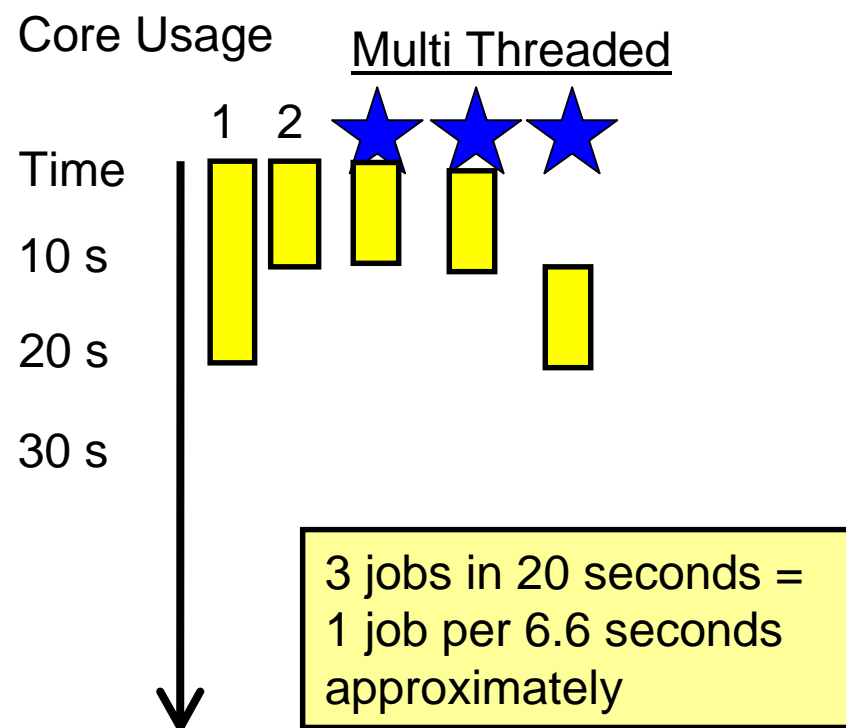

- Single CPU with requests taking 20 seconds to run sequentially – this consists of 10 seconds of CPU () and 10 seconds () waiting for disk (i.e. blocking)
- Three requests appear 10 seconds apart (indicated by )



Multiple CPUs/cores - increasing throughput

Consider this scenario:

- Dual-core CPU with requests taking 10 seconds to run sequentially – consisting entirely of CPU time (■)
- Three requests appear at the same time (indicated by ★)
- Ideal case - usually there are additional overheads.



Implementing a Multi-Threaded Server

- The multi-threaded File Reader Server **functionally** does exactly the same as the single-threaded File Reader Server – it sends back the contents of a specified file given a filename using an identical protocol.
- But its **non-functional** behaviour in terms of response times under load is very different.
- A key implementation difference is that it now starts a new Thread for every client request. This 'ServiceThread' actually carries out the required task.
- I have also handed out a 'TimedClient' listing which is identical to the previous listing except it also returns the time it takes for the server to respond.
- We will read through the code and discuss the changes.

Comparison between the single-threaded and multi-threaded servers

- We will run the single-threaded server with 3 simultaneous requests.
- Write down the response times taken here ...
- Based on the previous slides – what would you expect the 3 response times to be in the case of the multi-threaded server?
- Write down you predictions here ...

Are the response times as you predicted?

- We will now run the 3 simultaneous requests using the multi-threaded server ...
- Even though the tasks are not CPU bound we do not get the expected throughput if we assume the CPU is the main resource.
- Can you think of any reasons for this?
- How would you test these hypotheses?

Thread per Request Problems ...

So what's wrong with a Thread per request ?

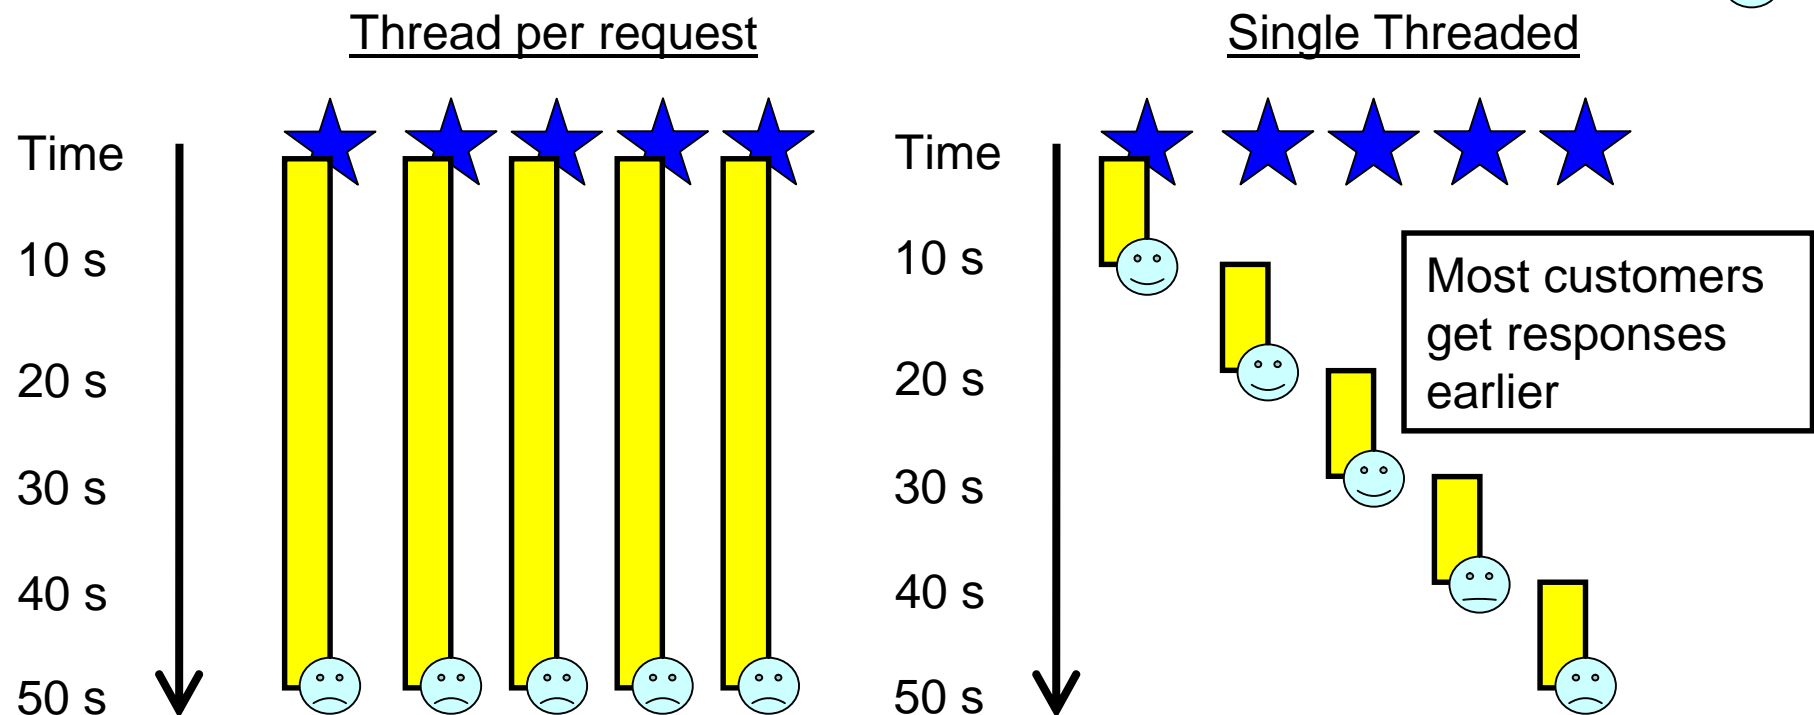
Not ideal (in some cases) because:

- As load on server increases the number of Threads created is unbounded ... resource allocation problems.
- Limited resources include CPU, memory, disk space, disk IO, network IO, OS limits on number of threads, etc.
- **Perceived throughput** is very poor for a heavily loaded system with a Thread per request.
- CPU-bound processes will get smaller and smaller time slices on the CPU as the number of Threads increase – leading to excessively long delays for ***most threads***.
- Similar situation for other machine resources ...

Perceived Throughput Scenario

Consider this scenario:

- Requests take 10 seconds of CPU and they are completely CPU bound (i.e. no I/O required).
- Five requests appear at about the same time – happy user?



A Thread Pool Helps Server Optimization

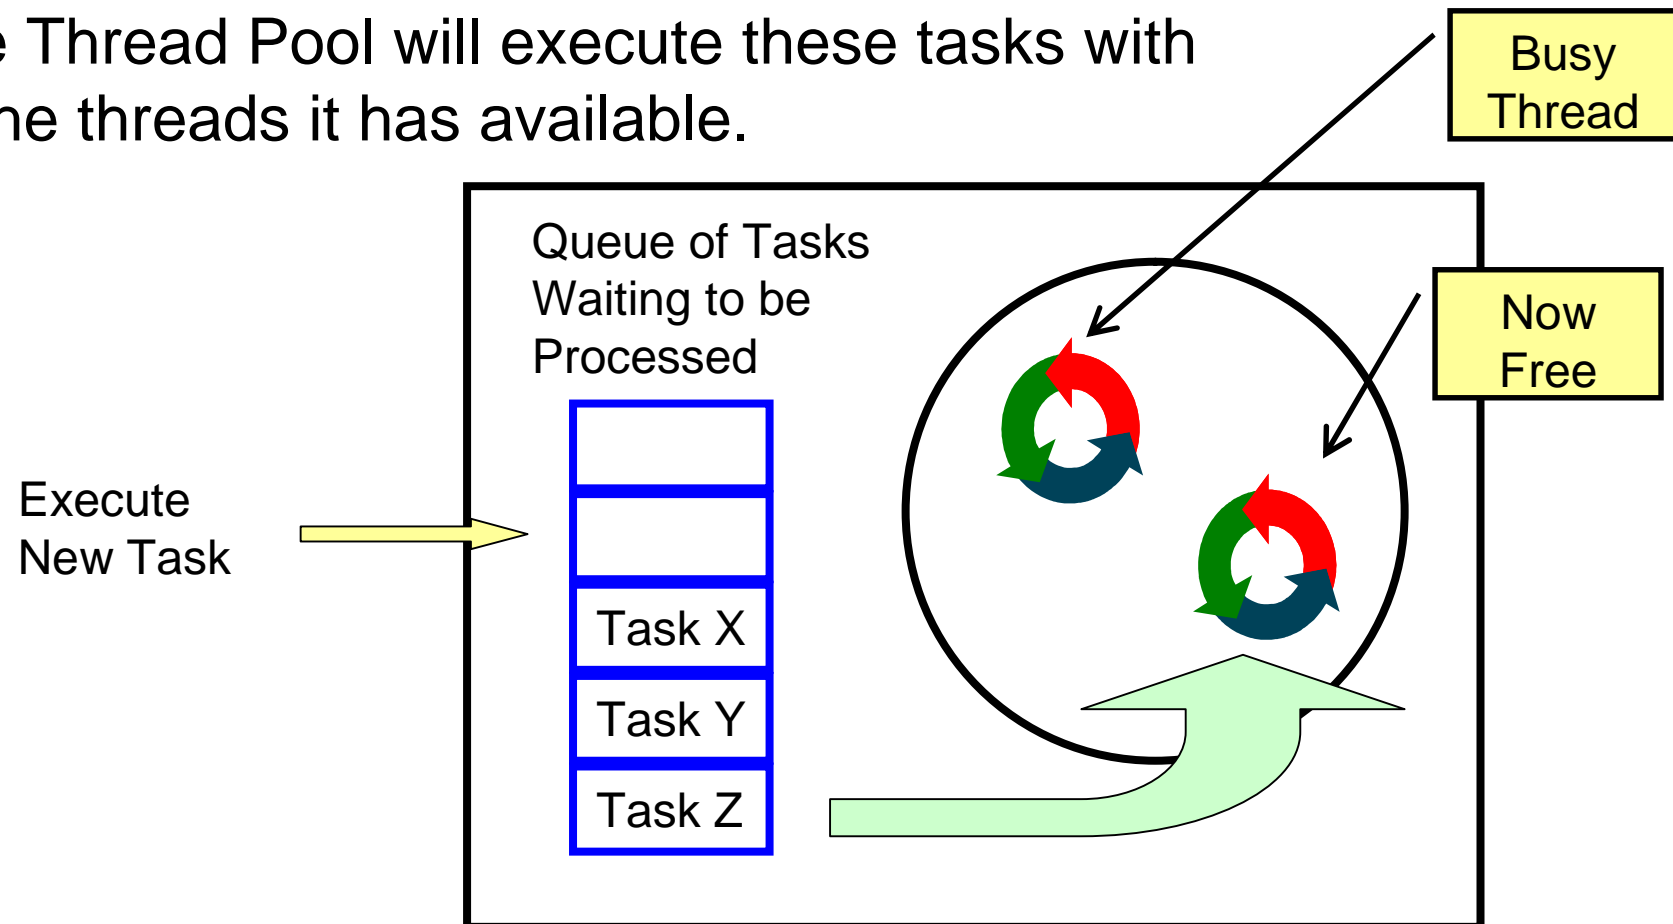
- Creating a Thread is not a trivial operation and uses CPU cycles. Each thread requires resources and having an unbounded number of Threads may result in resource exceptions such as OutOfMemory and inefficient use of resources.
- A Thread Pool consists of a pool of Threads that can be reused.
- The number of Threads within the pool can dynamically adjust **within set limits** to accommodate the current load.
- Thread Pool parameter optimization is complex and generally involves experimentation.

A Thread Pool Helps Server Optimization

Schematic of a Thread Pool with two Threads available

A set of tasks is assigned to the Thread Pool


The Thread Pool will execute these tasks with the threads it has available.



A Thread Pool allows optimization between different concerns such as CPU utilization, IO resources, throughput, etc.

Consists of a bounded number of available Threads (sometimes a fixed number) which are reused.

If we have a Pool of 2 Threads in the previous scenario.



Java 5 Executor Framework

- Fortunately Java 5 now has a standard framework for uncoupling the concerns about how an overall system creates 'tasks' (for instance `Runnable`s) and how it goes about executing these tasks.
- Classes that can execute 'tasks' implement the `Executor` interface.
- This allows a new type of `Executor` to be substituted into a system by only changing a single line of code (where it creates the concrete instance of the executor) ... rather than all the places that submit different types of 'tasks'.
- This 'task' decomposition of a concurrent system is a particularly powerful architecture.

```
package java.util.concurrent;  
  
public interface Executor {  
    void execute(Runnable command);  
}
```

Java 5 Thread Pool Implementation

- Within Java 5, a Thread Pool is a type of executor and can execute Runnable tasks.

```
Executor pool = new ThreadPoolExecutor(  
    10,      // Core Pool Size = 10 Threads.  
    20,      // Max Pool Size = 20 Threads.  
    50000L, // Thread idle 'timeout' = 50 seconds.  
    TimeUnit.MILLISECONDS,  
    new LinkedBlockingQueue<Runnable>());
```

- A task (implementing the interface Runnable) can be assigned to the pool using:

```
Runnable task = ... some Runnable object ...  
pool.execute(task);
```


Summary

- Optimization of multithreaded applications
 - Why a single threaded solution may not be ideal in some cases.
 - Why a thread-per-task solution may not be ideal in some cases.
 - The Thread Pool solution.
- A very brief introduction to the Java 5 Executor interface and Thread Pool class.